



#madmug

MongoDB + Hadoop

Norberto Leite

Wingman, MongoDB Inc.

Agenda

- MongoDB
- Hadoop
- MongoDB + Hadoop Connector
- How it works
- What can we do with it

MongoDB



MongoDB


The leading NoSQL database



General
Purpose




Document
Database




Open-
Source


Performance



Better Data
Locality



In-Memory Caching



In-Place
Updates

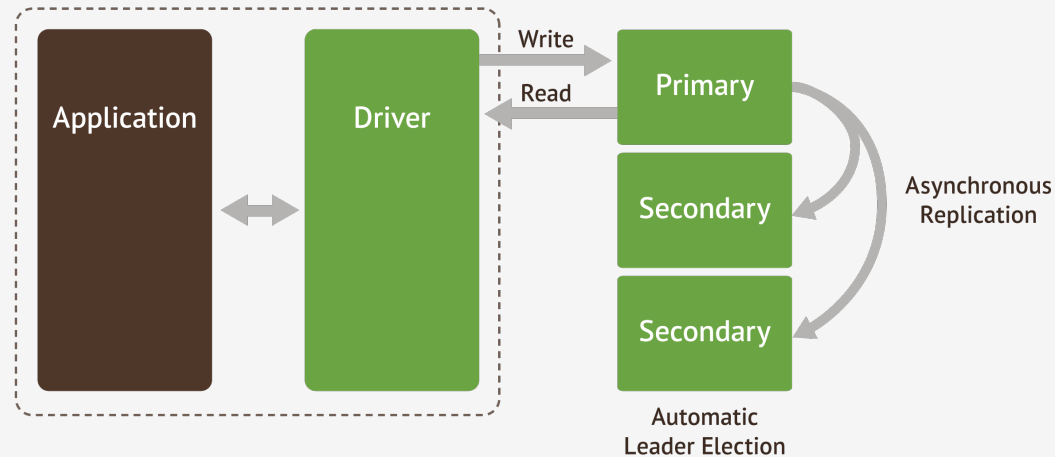
Scalability

Auto-Sharding



- Increase capacity as you go
- Commodity and cloud architectures
- Improved operational simplicity and cost visibility

High Availability



- Automated replication and failover
- Multi-data center support
- Improved operational simplicity (e.g., HW swaps)
- Data durability and consistency

Shell and Drivers

Drivers

Drivers for most popular programming languages and frameworks



Java



Ruby



JavaScript



Perl



Python



Haskell

Shell


Command-line shell for interacting directly with database

```
> db.collection.insert({product:"MongoDB",
type:"Document Database"})
>
> db.collection.findOne()
{
  "_id"       : ObjectId("5106c1c2fc629bfe52792e86"),
  "product"   : "MongoDB"
  "type"      : "Document Database"
}
```


Hadoop



Hadoop



- Google publishes seminal papers
 - GFS – global file system (Oct 2003)
 - MapReduce – divide and conquer (Dec 2004)
 - How they indexed the internet
- Yahoo builds and open sources (2006)
 - Doug Cutting lead, now at Cloudera
 - Most others now at Hortonworks
- Commonly mentioned has:
 - The elephant in the room!

Hadoop

- Primary components
 - HDFS – Hadoop Distributed File System
 - MapReduce – parallel processing engine
- Ecosystem
 - HIVE
 - HBASE
 - PIG
 - Oozie
 - Sqoop
 - Zookeeper





Apache Hadoop Ecosystem

Ambari

Provisioning, Managing and Monitoring Hadoop Clusters



Scoop

Data Exchange



Zookeeper

Coordination



Flume

Log Collector



Oozie

Workflow



Pig

Scripting



Mahout

Machine Learning

R Connectors

Statistics



Hive

SQL Query



Hbase

Columnar Store



HDFS

Hadoop Distributed File System

YARN Map Reduce v2

Distributed Processing Framework



MongoDB+Hadoop Connector

MongoDB+Hadoop Connector

CURRENT RELEASE: 1.0.0


The *Mongo+Hadoop Connector* (for brevity's sake, we'll refer to it as *mongo-hadoop* in this documentation) is a series of plugins for the [Apache Hadoop Platform](#) to allow connectivity to [MongoDB](#). This connectivity takes the form of allowing both reading MongoDB data into Hadoop (for use in MapReduce jobs as well as other components of the Hadoop ecosystem), as well as writing the results of Hadoop jobs out to MongoDB. A forthcoming release will also allow for reading and writing static BSON files (ala *mongodump* / *mongoexport*) to allow offline batching; commonly, users find this to be a beneficial feature to run analytics against backup data.

At this time, we support the "core" Hadoop APIs (now known as [Hadoop Common](#)), in the form of *mongo-hadoop-core*. There is additionally support for other pieces of the Hadoop Ecosystem, including [Pig](#) for ETL and [Streaming](#) for running Mongo+Hadoop jobs with Python (future releases will support additional scripting languages such as Ruby). Although it is not dependent upon Hadoop, we also provide a connector for the [Flume](#) distributed logging system.


MongoDB Hadoop Connector

- <http://api.mongodb.org/hadoop/MongoDB%2BHadoop+Connector.html>
- Allows interoperation of MongoDB and Hadoop
 - “Give power to the people”
- Allows processing across multiple sources
- Avoid custom hacky exports and imports
- Scalability and Flexibility to accommodate Hadoop and/or MongoDB changes


MongoDB with Hadoop



MongoDB with Hadoop



MongoDB with Hadoop



Benefits and Features

- Full multi-core parallelism to process data in MongoDB
- Full integration w/ Hadoop and JVM ecosystem
- Can be used on Amazon Elastic MapReduce
- Read and write backup files to local, HDFS and S3
- Vanilla Java MapReduce
 - But not only using Hadoop Streaming



Benefits and Features

- Supports PIG



- Supports HIVE



How it works

- Adapter examines MongoDB input collection and calculates a set of **splits** from data
- Each split is assigned to a Hadoop node
- In parallel hadoop pulls data from splits on MongoDB (or BSON) and starts processing locally
- Hadoop merges results and streams output back to MongoDB (or BSON) output collection

Example Tour



Tour bus stops

- Java MapReduce w/ MongoDB-Hadoop Connector
- Using Hadoop Streaming
- Pig and Hive

Data Set


- ENRON emails corpus (501 records, 1.75GB)
- Each document is one email
- <https://www.cs.cmu.edu/~enron/>

Document Example

```
{
  "_id" : ObjectId("4f2ad4c4d1e2d3f15a000000"),
  "body" : "Here is our forecast\n\n ",
  "filename" : "1.",
  "headers" : {
    "From" : "phillip.allen@enron.com",
    "Subject" : "Forecast Info",
    "X-bcc" : "",
    "To" : "tim.belden@enron.com",
    "X-Origin" : "Allen-P",
    "X-From" : "Phillip K Allen",
    "Date" : "Mon, 14 May 2001 16:39:00 -0700 (PDT)",
    "X-To" : "Tim Belden ",
    "Message-ID" : "<18782981.1075855378110.JavaMail.evans@thyme>",
    "Content-Type" : "text/plain; charset=us-ascii",
    "Mime-Version" : "1.0"
  }
}
```

Let's build a graph between
senders and recipients and
count the messages exchanged

Graph Sketch



Receiver Sender Pairs

```
{"_id": {"t": "bob@enron.com", "f": "alice@enron.com"}, "count" : 14}
{"_id": {"t": "bob@enron.com", "f": "eve@enron.com"}, "count" : 9}
{"_id": {"t": "alice@enron.com", "f": "charlie@enron.com"}, "count" : 99}
{"_id": {"t": "charlie@enron.com", "f": "bob@enron.com"}, "count" : 48}
{"_id": {"t": "eve@enron.com", "f": "charlie@enron.com"}, "count" : 20}
```

Vanilla Java MapReduce

Map Phase – each document get's through mapper function

@Override

```
public void map(NullWritable key, BSONObject val, final Context context){  
    BSONObject headers = (BSONObject)val.get("headers");  
    if(headers.containsKey("From") && headers.containsKey("To")){  
        String from = (String)headers.get("From"); String to = (String) headers.get("To");  
        String[] recips = to.split(",");  
        for(int i=0;i<recips.length;i++){  
            String recip = recips[i].trim();  
            context.write(new MailPair(from, recip), new IntWritable(1));  
        }  
    }  
}
```

Reduce Phase – output Maps are grouped by key and passed to Reducer

```
public void reduce( final MapPair pKey, final Iterable<IntWritable> pValues,
final Context pContext ){
    int sum = 0;
    for ( final IntWritable value : pValues ){
        sum += value.get();
    }
    BSONObject outDoc = new BasicDBObjectBuilder().start()
        .add( "f" , pKey.from)
        .add( "t" , pKey.to )
        .get();
    BSONWritable pkeyOut = new BSONWritable(outDoc);
    pContext.write( pkeyOut, new IntWritable(sum) ); }
```

Read From MongoDB (or BSON)

```
mongo.job.input.format=com.mongodb.hadoop.MongoInputFormat
```

```
mongo.input.uri=mongodb://my-db:27017/enron.messages
```

```
mongo.job.input.format=com.mongodb.hadoop.BSONFileInputFormat
```

```
mapred.input.dir= file:///tmp/messages.bson
```

```
mapred.input.dir= hdfs:///tmp/messages.bson
```

```
mapred.input.dir= s3:///tmp/messages.bson
```


Write To MongoDB (or BSON)

```
mongo.job.output.format=com.mongodb.hadoop.MongoOutputFormat
```

```
mongo.output.uri=mongodb://my-db:27017/enron.results_out
```

```
mongo.job.output.format=com.mongodb.hadoop.BSONFileOutputFormat
```

```
mapred.output.dir= file:///tmp/results.bson
```

```
mapred.output.dir= hdfs:///tmp/results.bson
```

```
mapred.output.dir= s3:///tmp/results.bson
```

Query Data


```
mongos> db.streaming.output.find({"_id.t": /^kenneth.lay/})
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "15126-1267@m2.innovyx.com" }, "count" : 1 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "2586207@www4.imakenews.com" }, "count" : 1 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "40enron@enron.com" }, "count" : 2 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "a..davis@enron.com" }, "count" : 2 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "a..hughes@enron.com" }, "count" : 4 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "a..lindholm@enron.com" }, "count" : 1 }
{ "_id" : { "t" : "kenneth.lay@enron.com",
            "f" : "a..schroeder@enron.com" }, "count" : 1 }
```

Hadoop Streaming



Same thing but lets use Python
PYTHON FTW!!!

Hadoop Streaming: How it works



Install python library

```
> pip install pymongo_hadoop
```

Map Phase

```
from pymongo_hadoop import BSONMapper
def mapper(documents): i=0
for doc in documents: i=i+1
    from_field = doc['headers']['From']
    to_field = doc['headers']['To']
    recips = [x.strip() for x in to_field.split(',')]
    for r in recips:
        yield {'_id': {'f':from_field, 't':r}, 'count': 1}
BSONMapper(mapper)
print >> sys.stderr, "Done Mapping."
```

Reduce Phase

```
from pymongo_hadoop import BSONReducer
def reducer(key, values):
    print >> sys.stderr, "Processing from/to %s" % str(key)
    _count = 0
    for v in values:
        _count += v['count']
    return {'_id': key, 'count': _count}
BSONReducer(reducer)
```


Surviving Hadoop



MapReduce easier w/ PIG and Hive

- PIG
 - Powerful language
 - Generates sophisticated map/reduce
 - Workflows from simple scripts
- HIVE
 - Similar to PIG
 - SQL as language

PIG



MongoDB+Hadoop and PIG

```
data = LOAD 'mongodb://localhost:27017/db.collection'  
        using com.mongodb.hadoop.pig.MongoLoader;
```

```
STORE records INTO 'file:///output.bson'  
        using com.mongodb.hadoop.pig.BSONStorage;
```

MongoDB + Hadoop and PIG

- PIG has a some special datatypes
 - Bags
 - Maps
 - Tuples
- MongoDB+Hadoop Connector converts between PIG and MongoDB datatypes

PIG

```
raw = LOAD 'hdfs:///messages.bson'
      using com.mongodb.hadoop.pig.BSONLoader(", 'headers:[]') ;
send_recip = FOREACH raw GENERATE $0#'From' as from, $0#'To'
as to;
//filter && split
send_recip_filtered = FILTER send_recip BY to IS NOT NULL;
send_recip_split = FOREACH send_recip_filtered GENERATE
      from as from, TRIM(FLATTEN(TOKENIZE(to))) as to;
//group && count
send_recip_grouped = GROUP send_recip_split BY (from, to);
send_recip_counted = FOREACH send_recip_grouped GENERATE
      group, COUNT($1) as count;
STORE send_recip_counted INTO 'file:///enron_results.bson' using
com.mongodb.hadoop.pig.BSONStorage;
```

HIVE



Hive w/ MongoDB-Hadoop

- Similar to Pig
 - Process data without need to write MapReduce from scratch
- But with SQL !

Data set example

```
db.users.find()
```

```
{ "_id": 1, "name": "Tom", "age": 28 }
```

```
{ "_id": 2, "name": "Alice", "age": 18 }
```

```
{ "_id": 3, "name": "Bob", "age": 29 }
```

```
{ "_id": 101, "name": "Scott", "age": 10 }
```

```
{ "_id": 104, "name": "Jesse", "age": 52 }
```

```
{ "_id": 110, "name": "Mike", "age": 32 }
```

```
...
```

Hive

```
CREATE TABLE mongo_users (id int, name string, age int)
STORED BY "com.mongodb.hadoop.hive.MongoStorageHandler"
WITH SERDEPROPERTIES( "mongo.columns.mapping" =
"_id,name,age" )
TBLPROPERTIES ( "mongo.uri" = "mongodb://localhost:27017/
test.users");
```

Hive

```
SELECT name,age FROM mongo_users WHERE id > 100 ;
```

```
SELECT * FROM mongo_users GROUP BY age WHERE id > 100 ;
```

```
SELECT * FROM mongo_users T1 JOIN user_emails T2 WHERE T1.id =  
T2.id;
```

Hive

```
INSERT OVERWRITE TABLE old_users SELECT id,name,age FROM  
mongo_users WHERE age > 100 ;
```

```
DROP TABLE mongo_users;
```

Upcoming



Roadmap Features

- Performance Improvements – Lazy BSON
- Full-Featured Hive Support
- Support multi-collection input
- API for **custom splitter** implementations
- And lots more ...

Recap

- Use Hadoop for massive MapReduce computations on big data sets stored in MongoDB
- MongoDB can be used as Hadoop filesystem
- There's lots of tools to make it easier
 - Streaming
 - Hive
 - PIG
 - EMR
- <https://github.com/mongodb/mongo-hadoop/tree/master/examples>



That's all Folks!

QA?





#madmug

MongoDB + Hadoop

Norberto Leite

Senior Solutions Architect, MongoDB Inc.