



Aggregation Framework

Agenda

- State of Aggregation
- Pipeline
- Usage and Limitations
- Optimization
- Sharding
- Expressions (time permitting)
- Looking Ahead

State of Aggregation



State of Aggregation

- We're storing our data in MongoDB
- We need to do ad-hoc reporting, grouping, common aggregations, etc.
- What are we using for this?

Data Warehousing



<http://www.kboil.net/warehouse1.jpg>

Data Warehousing

- SQL for reporting and analytics
- Infrastructure complications
 - Additional maintenance
 - Data duplication
 - ETL processes
 - Real time?

Aggregation Framework



<http://www.swissknifeshop.com/swiss-army-giant-by-wenger>

MapReduce

- Extremely versatile, powerful
- Intended for complex data analysis
- Overkill for simple aggregation tasks
 - Averages
 - Summation
 - Grouping

MapReduce in MongoDB

- Implemented with JavaScript
 - Single-threaded
 - Difficult to debug
- Concurrency
 - Appearance of parallelism
 - Write locks

Aggregation Framework



<http://www.victorinox.com/us/product/Swiss-Army-Knives/Category/Classics/Classic-SD/53001>

Aggregation Framework

- Declared in JSON, executes in C++
- Flexible, functional, and *simple*
 - Operation pipeline
 - Computational expressions
- Plays nice with sharding

Pipeline



Pipeline

- Process a stream of documents
 - Original input is a collection
 - Final output is a result document
- Series of operators
 - Filter or transform data
 - Input/output chain

Pipeline Operators

- \$match
- \$project
- \$group
- \$unwind
- \$sort
- \$limit
- \$skip

Our Example Data

```
{  
  _id: 375,  
  title: "The Great Gatsby",  
  ISBN: "9781857150193",  
  available: true,  
  pages: 218,  
  chapters: 9,  
  subjects: [  
    "Long Island",  
    "New York",  
    "1920s"  
  ],  
  language: "English"  
}
```

\$match

- Filter documents
- Uses existing query syntax
- No geospatial operations or \$where

Matching Field Values

```
{  
  title: "The Great Gatsby",  
  pages: 218,  
  language: "English"  
}
```



```
{ $match: {  
  language: "Russian"  
}}
```



```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

Matching with Query Operators

```
{  
  title: "The Great Gatsby",  
  pages: 218,  
  language: "English"  
}
```



```
{ $match: {  
  pages {$gt:100}  
}}
```



```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

\$project

- Reshape documents
- Include, exclude or rename fields
- Inject computed fields
- Create sub-document fields

Including and Excluding Fields

```
{
  _id: 375,
  title: "Great Gatsby",
  ISBN: "9781857150193",
  available: true,
  pages: 218,
  subjects: [
    "Long Island",
    "New York",
    "1920s"
  ],
  language: "English"
}
```



```
{ $project: {
  _id: 0,
  title: 1,
  language: 1
}}
```



```
{
  title: "Great Gatsby",
  language: "English"
}
```

Renaming and Computing Fields

```
{
  _id: 375,
  title: "Great Gatsby",
  ISBN: "9781857150193",
  available: true,
  pages: 218,
  chapters: 9,
  subjects: [
    "Long Island",
    "New York",
    "1920s"
  ],
  language: "English"
}
```



```
{ $project: {
  avgChapterLength: {
    $divide: ["$pages",
              "$chapters"]
  },
  lang: "$language"
}}
```



```
{
  _id: 375,
  avgChapterLength: 24.2222,
  lang: "English"
}
```

Creating Sub-Document Fields

```
{
  _id: 375,
  title: "Great Gatsby",
  ISBN: "9781857150193",
  available: true,
  pages: 218,
  chapters: 9,
  subjects: [
    "Long Island",
    "New York",
    "1920s"
  ],
  language: "English"
}
```



```
{ $project: {
  title: 1,
  stats: {
    pages: "$pages",
    language: "$language",
  }
}}
```



```
{
  _id: 375,
  title: "Great Gatsby",
  stats: {
    pages: 218,
    language: "English"
  }
}
```

\$group

- Group documents by an ID
 - Field reference, object, constant
- Other output fields are computed
 - \$max, \$min, \$avg, \$sum
 - \$addToSet, \$push
 - \$first, \$last
- Processes all data in memory

Calculating An Average

```
{  
  title: "The Great Gatsby",  
  pages: 218,  
  language: "English"  
}
```



```
{ $group: {  
  _id: "$language",  
  avgPages: { $avg:  
    "$pages" }  
}}
```



```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  _id: "Russian",  
  avgPages: 1440  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

```
{  
  _id: "English",  
  avgPages: 653  
}
```


Summating Fields and Counting

```
{  
  title: "The Great Gatsby",  
  pages: 218,  
  language: "English"  
}
```



```
{ $group: {  
  _id: "$language",  
  avgPages: { $avg:  
    "$pages" }  
}}
```



```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  _id: "Russian",  
  avgPages: 1440  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

```
{  
  _id: "English",  
  avgPages: 653  
}
```

Collecting Distinct Values

```
{  
  title: "The Great Gatsby",  
  pages: 218,  
  language: "English"  
}
```



```
{ $group: {  
  _id: "$language",  
  titles: { $addToSet:  
    "$title" }  
}}
```



```
{  
  title: "War and Peace",  
  pages: 1440,  
  language: "Russian"  
}
```

```
{  
  _id: "Russian",  
  titles: ["War and Peace"]  
}
```

```
{  
  title: "Atlas Shrugged",  
  pages: 1088,  
  language: "English"  
}
```

```
{  
  _id: "English",  
  titles: [  
    "Atlas Shrugged",  
    "The Great Gatsby" ]  
}
```

\$unwind

- Operate on an array field
- Yield new documents for each array element
 - Array replaced by element value
 - Missing/empty fields → no output
 - Non-array fields → error
- Pipe to \$group to aggregate array values

Collecting Distinct Values

```
{  
  title: "The Great  
Gatsby",  
  ISBN: "9781857150193",  
  subjects: [  
    "Long Island",  
    "New York",  
    "1920s"  
  ]  
}
```



```
{ $unwind: "$subjects" }
```



```
{ title: "The Great Gatsby",  
  ISBN: "9781857150193",  
  subjects: "Long Island" }
```

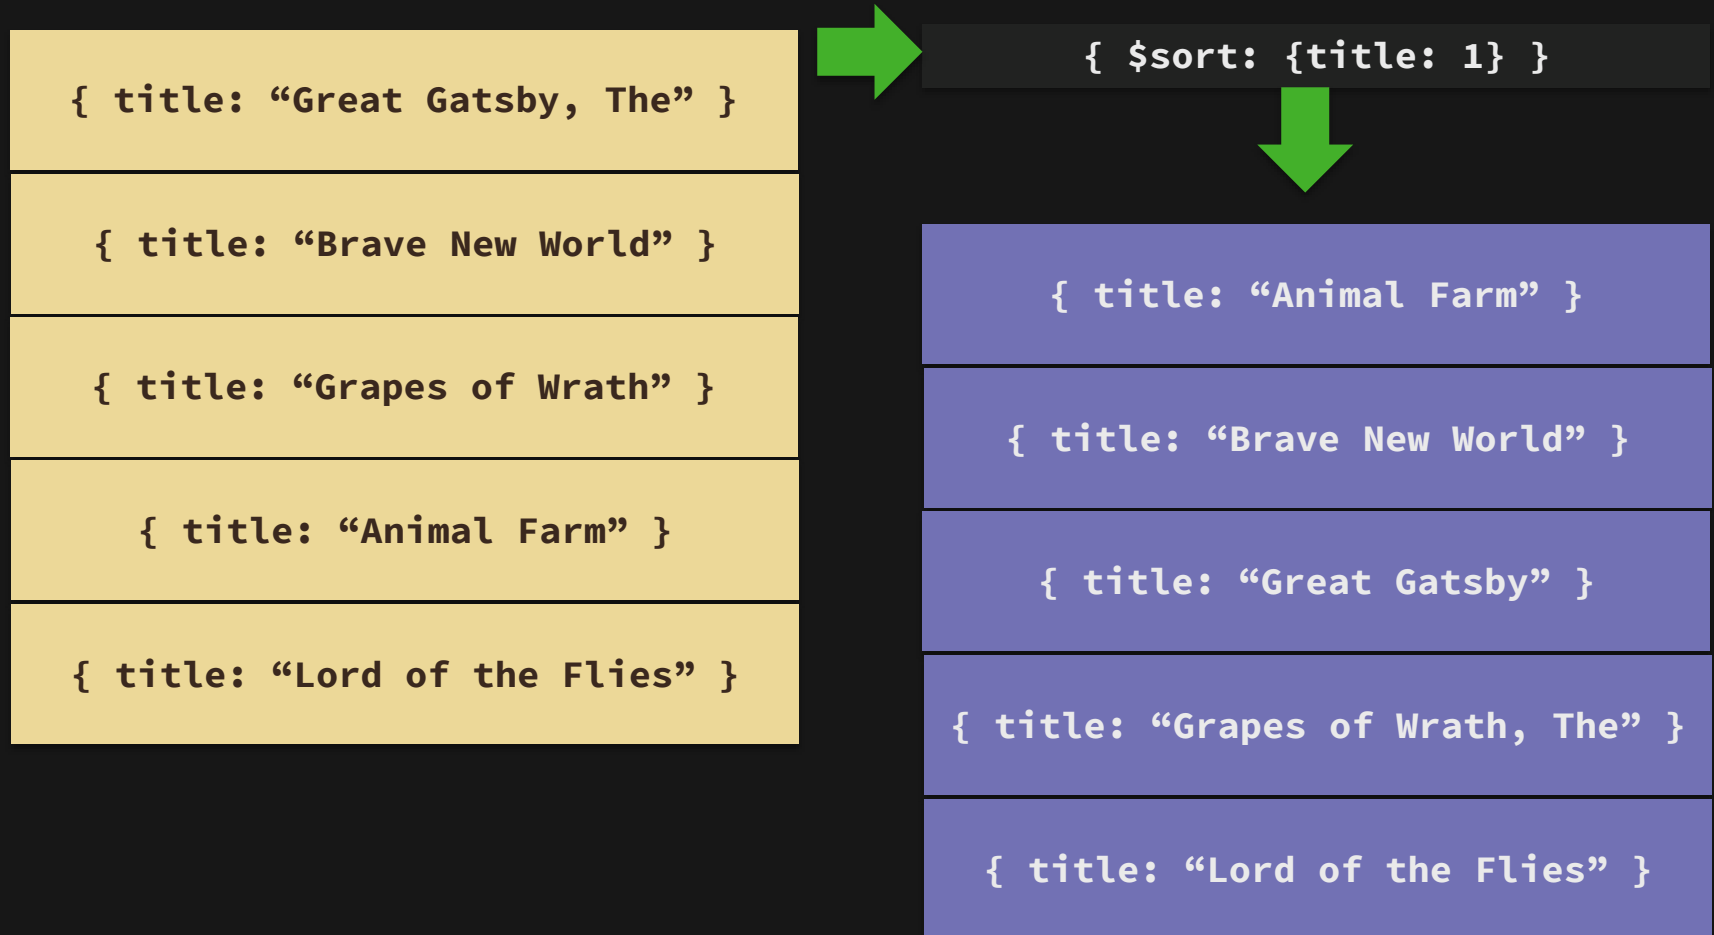
```
{ title: "The Great Gatsby",  
  ISBN: "9781857150193",  
  subjects: "New York" }
```

```
{ title: "The Great Gatsby",  
  ISBN: "9781857150193",  
  subjects: "1920s" }
```

\$sort, \$limit, \$skip

- Sort documents by one or more fields
 - Same order syntax as cursors
 - Waits for earlier pipeline operator to return
 - In-memory unless early and indexed
- Limit and skip follow cursor behavior

Sort All the Documents in the Pipeline



Limit Documents Through the Pipeline



Limit Documents Through the Pipeline



Usage and Limitations

Usage

- `collection.aggregate()` method
 - Mongo shell
 - Most drivers
- `aggregate` database command

Collection

```
db.books.aggregate([
  { $project: { language: 1 } },
  { $group: { _id: "$language", numTitles: { $sum: 1 } } }
])
```



```
{
  result: [
    { _id: "Russian", numTitles: 1 },
    { _id: "English", numTitles: 2 }
  ],
  ok: 1
}
```

Database Command

```
db.runCommand({
  aggregate: "books",
  pipeline: [
    { $project: { language: 1 } },
    { $group: { _id: "$language", numTitles: { $sum: 1 } } }
  ]
})
```



```
{
  result: [
    { _id: "Russian", numTitles: 1 },
    { _id: "English", numTitles: 2 }
  ],
  ok: 1
}
```

Limitations

- Result limited by BSON document size
 - Final command result
 - Intermediate shard results
 - **cursor and \$out variants in MongoDB 2.5 beta!**
- Pipeline operator memory limits
- Some BSON types unsupported
 - Binary, Code, deprecated types

Sharding



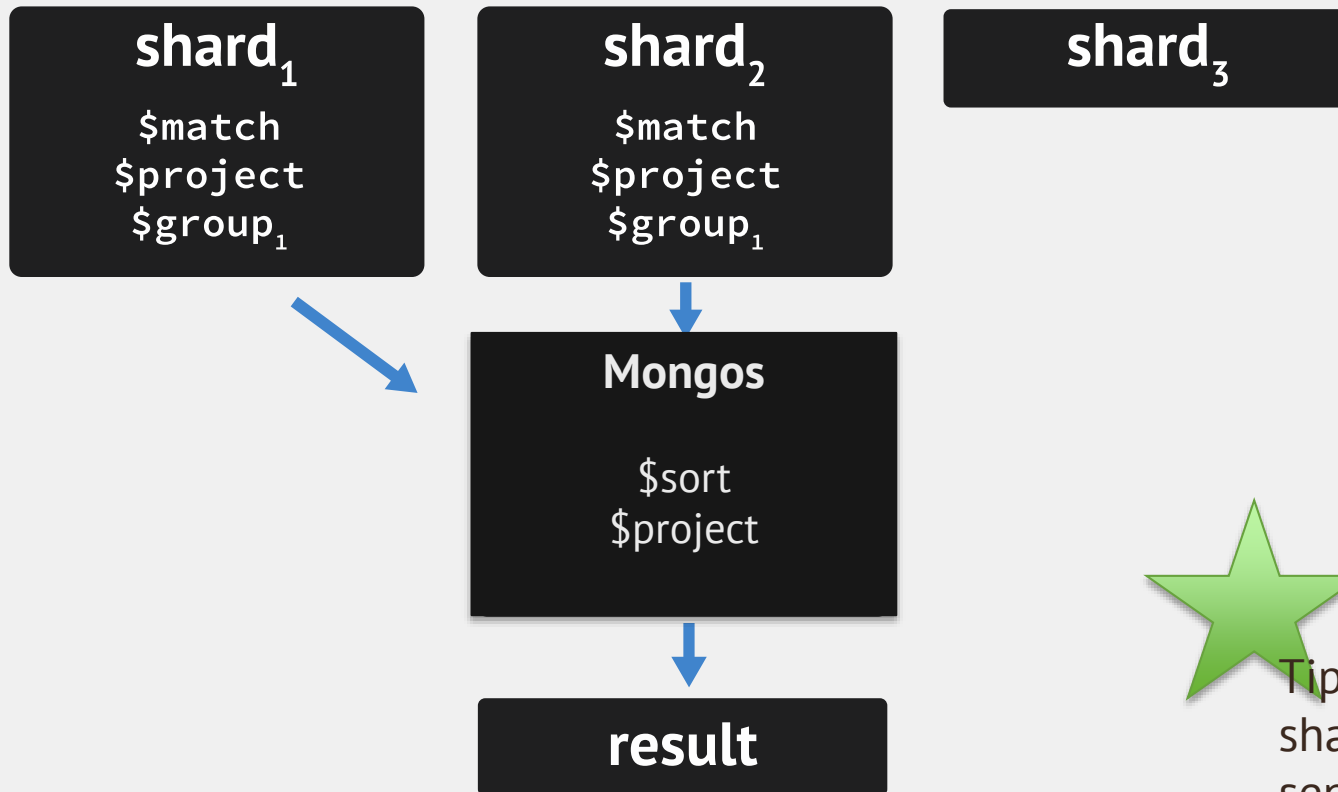
Sharding

- Split the pipeline at first \$group or \$sort
 - Shards execute pipeline up to that point
 - mongos merges results and continues
- Early \$match may excuse shards
- CPU and memory implications for mongos

Sharding

```
[  
  { $match:    { /* filter by shard key */ }},  
  { $project: { /* select fields          */ }},  
  { $group:    { /* group by some field  */ }},  
  { $sort:     { /* sort by some field   */ }},  
  { $project: { /* reshape result        */ }}  
]
```


Sharding



Tip: Use profiler on each shard to see what mongos sent to the shard.

Looking Ahead



Extending the Framework

- Adding new pipeline operators, expressions
- Removing 16MB limit for result set
- \$out and “tee” for output control
 - <https://jira.mongodb.org/browse/SERVER-3253>
- explain()



Thank You

