




Indexing and Query Optimization

Agenda

- What are indexes?
- Why do I need them?
- Working with indexes in MongoDB
- Optimize your queries
- Avoiding common mistakes

What Are Indexes?



<http://homemaker2words.com/>

What Are Indexes?

Imagine you're looking for a recipe in a cookbook ordered by recipe name. Looking up a recipe by name is quick and easy.


Consult the Index

Z A B C D E F G H I J K L M N O P Q R S T U V W X Y Z	SUPER EASY BURRITOS FOR A CROWD	118	SAUSAGE BISCUITS	139	EASY CHERRY CHEESE CAKE	161	ZUCCHINI CUSTARD PIE	180
	SWEET AND SOUR CHICKEN WITH RICE	118	STICKY BUN FRENCH TOAST	140	EGG CUSTARD PIE	161	COOKIES & CANDY	
	SWEET CAJUN BEEF STEW	119	STICKY BUNS	140	EGGNOG PIE	162	AUNT FLORENCE'S CHRISTMAS COOKIES	181
	SWEET N SOUR CHICKEN BREASTS	120	SUNFLOWER BREAD	141	FLAT APPLE PIE	162	BEST EVER CHOCOLATE CHIP COOKIES	181
	TAMALES	120	SUPER EASY INDIAN CHAPATIS	141	FLOURLESS CHOCOLATE CAKE	163	BISCOTTI QUEEN'S ITALIAN LEMON BISCOTTI	182
	THE KING'S QUICHE	122	SWEET CORNBREAD CAKE	142	FOOLPROOF FLAN	163	CHOCOLATE CHIP COOKIE BARS	183
	TOUCAN CHILI	122	WHITE BREAD (GUMPBREAD)	142	FRESH STRAWBERRY CAKE	164	CHOCOLATE CRACKER CRUNCH BARS	183
	WHITE CHICKEN CHILI (CROCK POT RECIPE)	123	WHOLE WHEAT BANANA BREAD	143	GOLDEN SPONGE CAKE	164	CHOCOLATE RIBBONS	183
	ZESTY CHEESY RAVIOLI	124	ZUCCHINI BREAD	144, 145	GRANDMA BURNEY'S BLUEBERRY CRUNCH	165	CREAMY PRALINES	184
	BREADS & ROLLS		DESSERTS		GRANDMA SADIE'S CREAMY CHEESECAKE	165	EASY FUDGE	184
	BANANA CREAM CHEESE BREAD	125	10 EGG POUND CAKE	147	GRANDMA'S EGG CUSTARD	166	FAST PEANUT BUTTER FUDGE	185
	BEER CHEESE BREAD	125	216 AND ¾ PIE CRUST	147	HOMEMADE ICE CREAM	166	FORGET ME COOKIES	185
	BEST BANANA BREAD EVER	125	5 MINUTE CHOCOLATE MUG CAKE	148	LEMON CHESS PIE	166	FRUITY SPRITZ COOKIES	185
	BRAN OVERNIGHT REFRIGERATOR ROLLS	126	AMISH APPLE PIE	148	LOVELY LEMON PIE	167	GRAHAM CRACKER MERINGUE COOKIES	186
	BROCCOLI CORNBREAD	126	ANGEL FOOD CAKE	149	MAMAW'S CHOCOLATE PIE	167	GRAMA PAT'S OATMEAL COOKIES	186
	BUTTER DIPS	127	APPLE CINNAMON CAKE	149	MAW'S PECAN PIE	168	GRANDMA FLO'S SUPER EASY ROCKY ROAD	187
	BUTTER HORN ROLLS	127	APPLE CRISP	150	MOIST LEMON LOVER'S POUND CAKE	168	GRANDMA'S CRUNCHY COOKIES	187
	CARROT COCONUT BREAD	128	APPLE SPICE CAKE WITH CREAM CHEESE FROSTING	150	MOUNTAIN MAMMA MUD PUDDING CAKE	169	HOLIDAY SUGAR COOKIES	187
	CHEESE BISCUITS	128	AUNT SALLY'S BUÑUELOS	151	NANNIES' EGG PIE	169	LACE COOKIES	188
	CHEROKEE FRYBREAD	128	BAKED CUSTARD	151	NO-BAKE CHEESECAKE	170	MICROWAVE CARMEL POPCORN	188
	CHERRY CORNBREAD	129	BAKLAVA	151	OATMEAL CHOCOLATE BARS	170	MOCHA FROSTED DROPS	189
	CINNAMON SWIRL BREAD	129	BERRY TRIFLE	152	OLD FASHIONED POUND CAKE	171	MOLASSES COOKIES	189
	CRANBERRY BREAD	130	BLACKBERRY COBBLER	153	ORANGE BALLS	171	MOMMA'S SUGAR COOKIES	190
	EASY SOFT RYE BREAD	131	BLUEBERRY COOLWHIP PIE	153	ORANGE CRANBERRY BARS	172	MOM'S SUGAR COOKIES	190
	EGG BREAD	131	BREAD AND BUTTER PUDDING	153	PEACH COBBLER	172	NO-BAKE COOKIES	191
	FRENCH TOAST CUSTARD	132	BUTTER POUND CAKE	154	PEACH CRISP	172	NO-BAKE COOKIES	191
	GOLDEN CRESCENT ROLLS	132	BUTTERSCOTCH PIE	154	PEACH PIZZA	173	OATMEAL PEANUT BUTTER CHOCOLATE CHIP	191
	GRAMA'S FLOUR TORTILLAS	133	CAKE & ICE CREAM DELIGHT	154	PEANUT BUTTER ICING	173		
	KUCHEN BREAD (GERMAN SWEET BREAD)	133	CARAMEL FLAN	155	PERFECT CHERRY PIE	173		
	LIGHTER THAN AIR PAN ROLLS	134	CHEESECAKE WITH PASTRY	156	PUMPKIN DELIGHT	174		
	http://files.backyardchickens.com/images/BYC-Cookbook-index-3.gif				PUMPKIN DESSERT	174		


Linked List



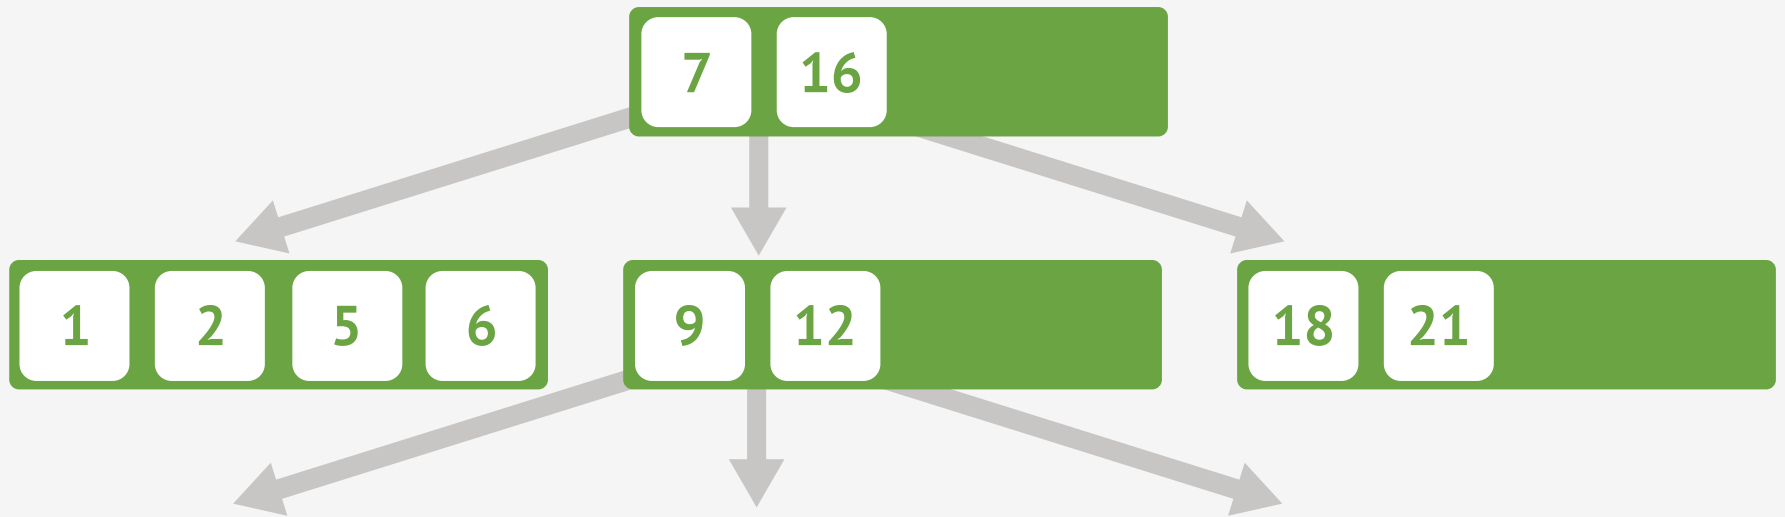
Finding 7 in a Linked List



Finding 7 In a Tree




Indexes in MongoDB are B-Trees



Queries, inserts and deletes:
 $O(\log(n))$ time

B-Trees



Indexes are the single biggest tunable performance factor in MongoDB

Absent or suboptimal indexes
are the most common
avoidable MongoDB
performance problem.

Why do I need indexes?

A brief story

Working with Indexes in MongoDB

How Do I Create Indexes?

```
// Create an index if one does not exist  
db.recipes.ensureIndex ( { main_ingredient : 1 } )
```

```
// The client remembers the index and raises no errors  
db.recipes.ensureIndex( { main_ingredient : 1 } )
```

* 1 means ascending, -1 descending

What Can Be Indexed?

// Multiple fields (compound key indexes)

```
db.recipes.ensureIndex({  
  main_ingredient: 1,  
  calories: -1  
})
```

// Arrays of values (multikey indexes)

```
{  
  name: 'Chicken Noodle Soup',  
  ingredients : [ 'chicken', 'noodles' ]  
}
```

```
db.recipes.ensureIndex( { ingredients : 1 } )
```

What Can Be Indexed?

```
// Embedded documents
```

```
{  
  name : 'Apple Pie',  
  contributor : {  
    name : 'Joe American',  
    id : 'joea123'  
  }  
}
```

```
db.recipes.ensureIndex ( { 'contributor.id' : 1 } )
```

```
db.recipes.ensureIndex ( { 'contributor' : 1 } )
```

How Do I Manage Indexes?

// List a collection's indexes

```
db.recipes.getIndexes()
```

```
db.recipes.getIndexKeys()
```

// Drop a specific index

```
db.recipes.dropIndex ( { ingredients : 1 } )
```

// Drop all indexes and recreate them

```
db.recipes.reIndex( )
```

// Default (unique) index on _id

Background Index Builds

// Index creation is a blocking operation that can take a long time

// Background creation yields to other operations

```
db.recipes.ensureIndex (
```

```
  { ingredients: 1 },
```

```
  { background: true }

```

```
)
```

Options

- Uniqueness constraints (unique, dropDups)
- Sparse Indexes

Uniqueness Constraints

// Only one recipe can have a given value for name

```
db.recipes.ensureIndex( { name: 1 }, { unique: true } )
```

// Force index on collection with duplicate recipe names – drop the duplicates

```
db.recipes.ensureIndex(  
  { name: 1 },  
  { unique: true, dropDups: true }  
)
```

* dropDups should be used with caution

Sparse Indexes

// Only documents with field calories will be indexed

```
db.recipes.ensureIndex(
```

```
  { calories : -1 },
```

```
  { sparse : true }
```

```
)
```

// Allow multiple documents to not have calories field

```
db.recipes.ensureIndex (
```

```
  { name: 1 , calories : -1 },
```

```
  { unique : true, sparse : true }
```

```
)
```

* Missing fields are stored as null(s) in the index

Other Index Types

- Geospatial Indexes (2d Sphere)
- Text Indexes
- TTL Collections (expireAfterSeconds)
- Hashed Indexes for sharding

Geospatial Indexes

```
// Add GeoJSON
```

```
{
```

```
  name: 'MongoDB Palo Alto',
```

```
  loc: { type : "Point",
```

```
    coordinates: [ 37.449157 , -122.158574 ] }
```

```
}
```

```
// Index the coordinates
```

```
db.locations.ensureIndex( { loc : '2dsphere' } )
```

Geospatial Indexes

// Query for locations 'near' a particular coordinate

```
db.locations.find ( {  
  loc: { $near:  
    { $geometry:  
      { type : "Point",  
        coordinates: [ 37.4, -122.3 ] },  
      $maxDistance: 40 }  
    }  
  } )
```

Text Indexes

```
db.recipes.insert( { _id : 4 , y : "add flour and mix" } );  
db.recipes.ensureIndex( { y : "text" } );
```

Limitations

- Collections can not have > 64 indexes.
- Index keys can not be > 1024 bytes (1K).
- The name of an index, including the namespace, must be < 125 characters.
- Queries can only use 1 index*
- Indexes have storage requirements, and impact the performance of writes.
- In memory sort (no-index) limited to 32mb of return data.

Optimize Your Queries

The “Explain” Plan (Pre-Index)

```
db.recipes.find( { calories:
  { $lt : 40 } }
).explain( )
{
  "cursor" : "BasicCursor" ,
  "n" : 42,
  "nscannedObjects" : 12345,
  "nscanned" : 12345,
  ...
  "millis" : 356,
  ...
}
```

* Doesn't use cached plans, re-evals and resets cache

The “Explain” Plan (Post-Index)

```
db.recipes.find( { calories:
  { $lt : 40 } }
).explain( )
{
  "cursor" : "BtreeCursor calories_1" ,
  "n" : 42,
  "nscannedObjects" : 42,
  "nscanned" : 42,
  ...
  "millis" : 3,
  ...
}
```

* Doesn't use cached plans, re-evals and resets cache

Profiling Slow Operations

```
db.setProfilingLevel( n , slowms=100ms )
```

n=0 profiler off

n=1 record operations longer than *slowms*

n=2 record all queries

```
db.system.profile.find()
```

*** The profile collection is a capped collection, and fixed in size**

The Query Optimizer

- For each "type" of query, MongoDB periodically tries *all* useful indexes
- Aborts the rest as soon as one plan wins
- The winning plan is temporarily cached for each "type" of query

Manually Select Index to Use

// Tell the database what index to use

```
db.recipes.find( {  
  calories: { $lt : 1000 } }  
).hint ( { _id : 1 } )
```

// Tell the database to NOT use an index

```
db.recipes.find(  
  { calories: { $lt : 1000 } }  
).hint( { $natural : 1 } )
```

Use Indexes to Sort Query Results

// Given the following index

```
db.collection.ensureIndex( { a : 1, b : 1 , c : 1, d : 1 } )
```

// The following query and sort operations can use the index

```
db.collection.find( ).sort( { a : 1 } )
```

```
db.collection.find( ).sort( { a : 1, b : 1 } )
```

```
db.collection.find( { a : 4 } ).sort( { a : 1, b : 1 } )
```

//This query will not use the index, but will use the index to sort

```
db.collection.find( { b : 5 } ).sort( { a : 1, b : 1 } )
```

Indexes that won't work for sorting query results

// Given the following index

```
db.collection.ensureIndex( { a : 1, b : 1, c : 1, d : 1 } )
```

// These can not sort using the index

```
db.collection.find( ).sort( { b : 1 } )
```

```
db.collection.find( { b : 5 } ).sort( { b : 1 } )
```

Indexes that won't work for sorting query results

```
// MongoDB can return data from just the index
db.recipes.ensureIndex({ main_ingredient : 1, name : 1 })
```

```
// Return only the ingredients field
db.recipes.find(
  { main_ingredient : 'chicken' },
  { _id: 0, name : 1 }
)
```

```
// indexOnly will be true in the explain plan
db.recipes.find(
  { main_ingredient : 'chicken' },
  { _id: 0, name: 1 }
).explain()
{
  "indexOnly" : true,
}
```


Absent or suboptimal indexes
are the most common
avoidable MongoDB
performance problem.

Avoiding Common Mistakes

Trying to Use Multiple Indexes

// MongoDB can only use one index for a query

```
db.collection.ensureIndex( { a : 1 } )
```

```
db.collection.ensureIndex( { b : 1 } )
```

// Only one of the above indexes is used

```
db.collection.find( { a : 3, b : 4 } )
```

Compound Key Mistakes

// MongoDB can only use one index for a query

```
db.collection.ensureIndex({ a: 1 })
```

```
db.collection.ensureIndex({ b: 1 })
```

// Only one of the above indexes is used

```
db.collection.find({ a: 3, b: 4 })
```

// Use a compound index

```
db.collection.ensureIndex( { a : 1, b : 1 } )
```

Low Selectivity Indexes

```
db.collection.distinct('status')  
[ 'new', 'processed' ]
```

```
db.collection.ensureIndex( { status : 1 } )
```

// Low selectivity indexes provide little benefit

```
db.collection.find( { status : 'new' } )
```

// Better

```
db.collection.ensureIndex ( { status : 1, created_at : -1 } )
```

```
db.collection.find(  
  { status : 'new' }  
) .sort( { created_at : -1 } )
```

Regular Expressions

```
db.users.ensureIndex( { username : 1 } )
```

```
// Left anchored regex queries can use the index
```

```
db.users.find( { username : /^joe smith/ } )
```

```
// But not generic regexes
```

```
db.users.find( {username : /smith/ } )
```

```
// Or case insensitive queries
```

```
db.users.find( { username : /Joe/i } )
```

Negation

// Indexes aren't helpful with negations

```
db.things.ensureIndex( { x : 1 } )
```

// e.g. "not equal" queries

```
db.things.find( { x : { $ne : 3 } } )
```

// ...or "not in" queries

```
db.things.find( { x : { $nin : [2, 3, 4] } } )
```

// ...or the \$not operator

```
db.people.find( { name: { $not : 'John Doe' } } )
```

Choosing the right indexes is one of the most important things you can do as a MongoDB developer.
Take indexes into consideration when designing your application.



Thank You

