



Schema Design





Agenda

- Working with documents
- Evolving a Schema
- Queries and Indexes
- Common Patterns

RDBMS		MongoDB
Database	→	Database
Table	→	Collection
Row	→	Document
Index	→	Index
Join	→	Embedded Document
Foreign Key	→	Reference

Working with Documents

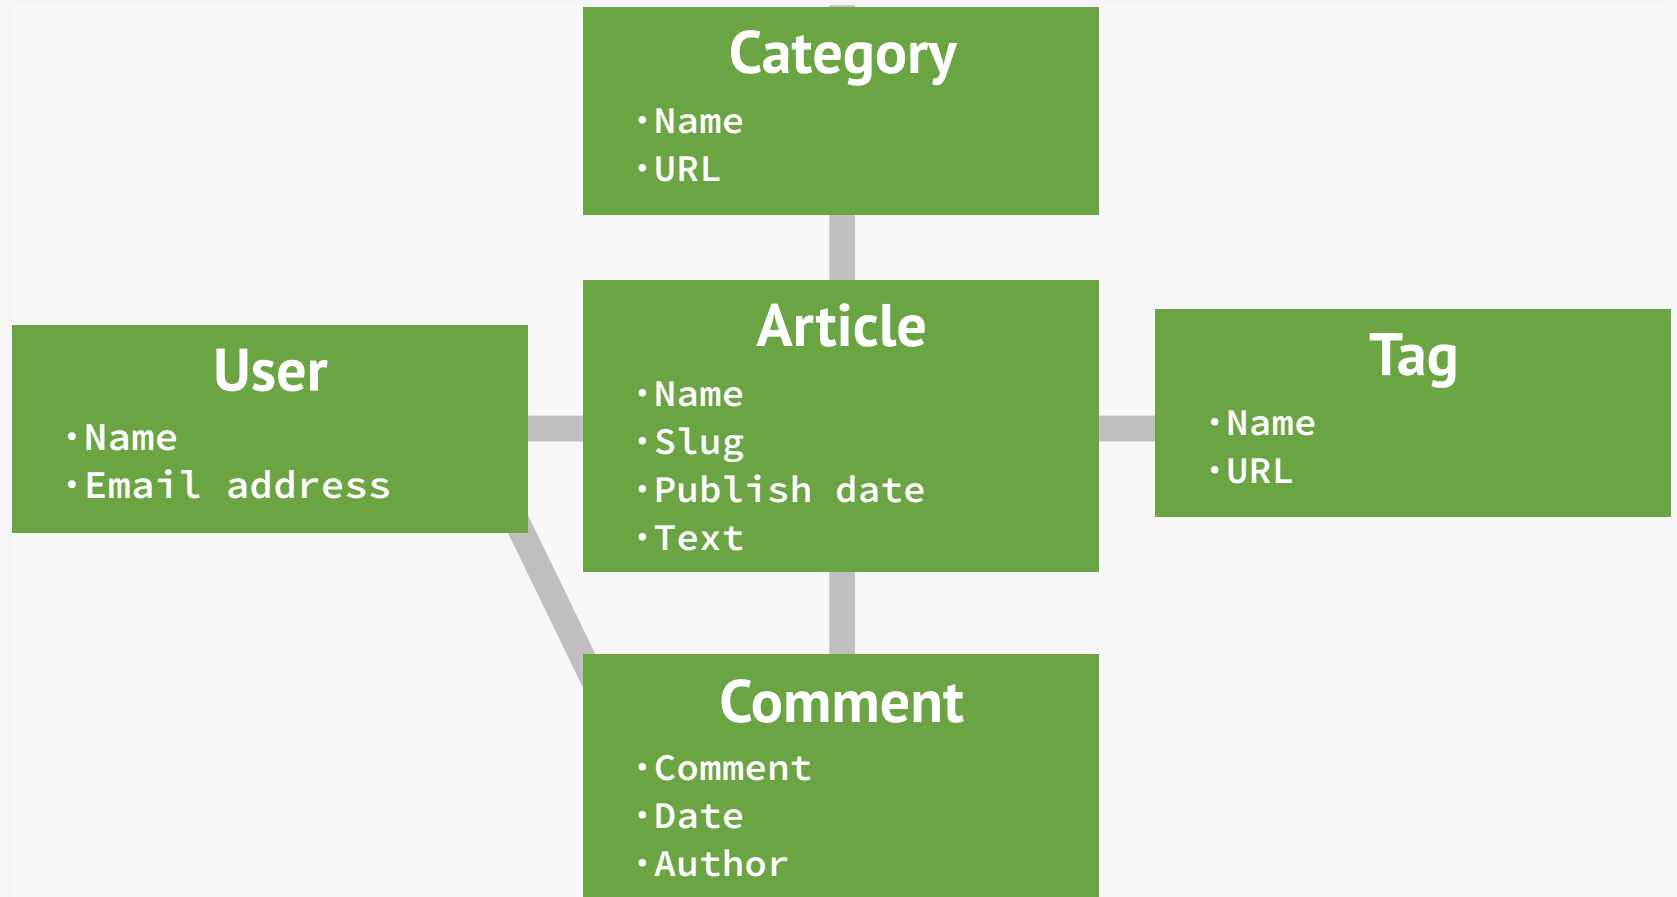
Modeling Data



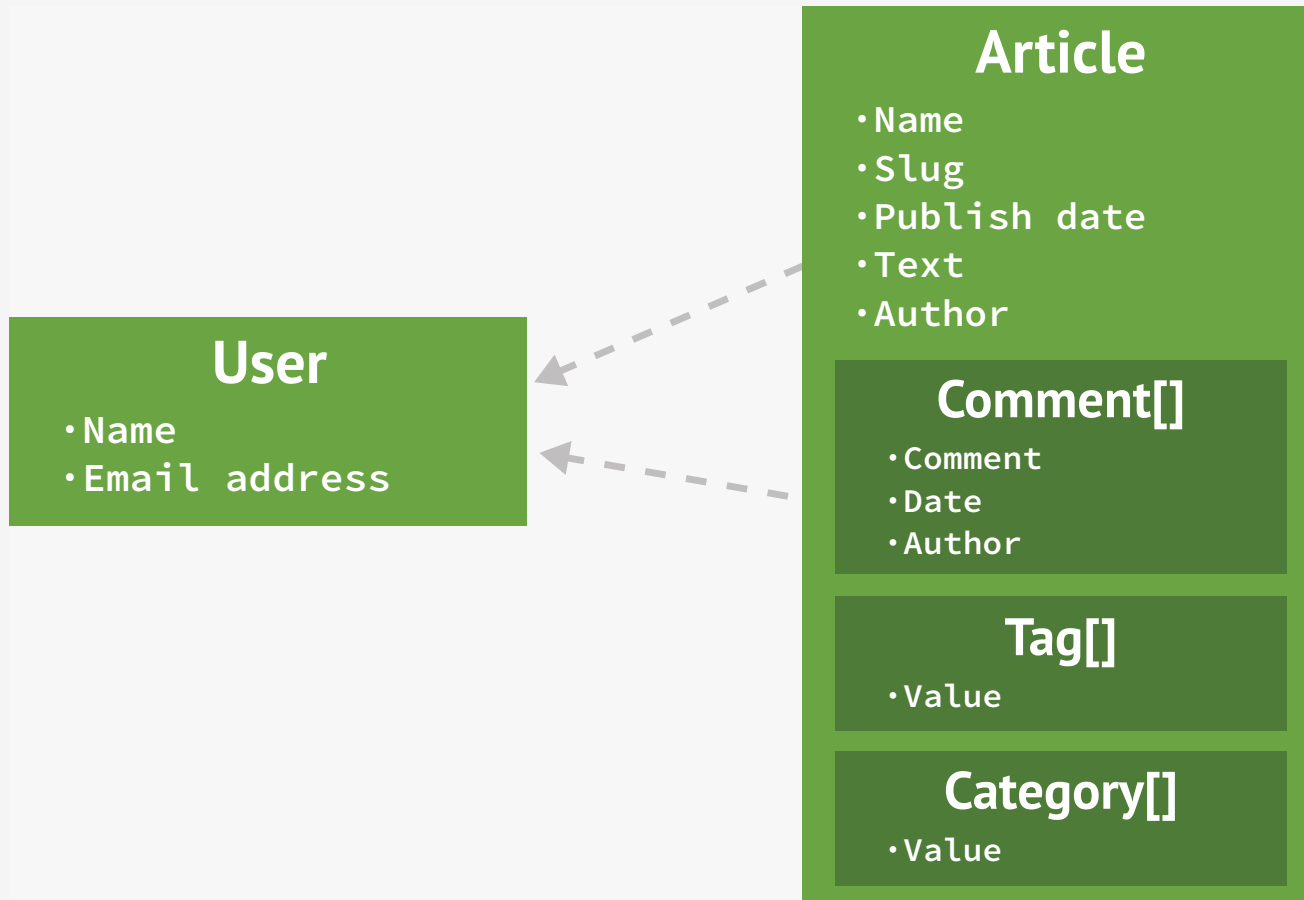
Documents

**Provide flexibility and
performance**

Normalized Data



De-Normalized (embedded) Data



Relational Schema Design

Focus on data storage

Document Schema Design

Focus on data use

Schema Design Considerations

- How do we manipulate the data?
 - Dynamic Ad-Hoc Queries
 - Atomic Updates
 - Map Reduce
- What are the access patterns of the application?
 - Read/Write Ratio
 - Types of Queries / Updates
 - Data life-cycle and growth rate

Data Manipulation

Query Selectors

- **Scalar:** \$ne, \$mod, \$exists, \$type, \$lt, \$lte, \$gt, \$gte
- **Vector:** \$in, \$nin, \$all, \$size

Atomic Update Operators

- **Scalar:** \$inc, \$set, \$unset
- **Vector:** \$push, \$pop, \$pull, \$pushAll, \$pullAll, \$addToSet

Data Access

Flexible Schemas

Ability to embed complex data structures

Secondary Indexes

Multi-Key Indexes

Aggregation Framework

- \$project, \$match, \$limit, \$skip, \$sort, \$group, \$unwind

No Joins

Getting Started



Library Management Application

Patrons

Books

Authors

Publishers

An Example

One to One Relations

Modeling Patrons

```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader"  
}
```

```
address = {  
  patron_id = "joe",  
  street: "123 Fake St. ",  
  city: "Faketon",  
  state: "MA",  
  zip: 12345  
}
```



```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader",  
  address: {  
    street: "123 Fake St. ",  
    city: "Faketon",  
    state: "MA",  
    zip: 12345  
  }  
}
```

One to One Relations

Mostly the same as the relational approach

Generally good idea to embed “contains” relationships

Document model provides a holistic representation of objects

An Example

One To Many Relations

Modeling Patrons

```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader",  
  join_date: ISODate("2011-10-15"),  
  addresses: [  
    {street: "1 Vernon St.", city: "Newton", state: "MA", ...},  
    {street: "52 Main St.", city: "Boston", state: "MA", ...}  
  ]  
}
```


Publishers and Books

- Publishers put out many books
- Books have one publisher

Book

MongoDB: The Definitive Guide,
By Kristina Chodorow and Mike Dirolf
Published: 9/24/2010
Pages: 216
Language: English

Publisher: O'Reilly Media, CA



Modeling Books – Embedded Publisher

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher: {  
    name: "O'Reilly Media",  
    founded: "1980",  
    location: "CA"  
  }  
}
```

Modeling Books & Publisher Relationship

```
publisher = {  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
}
```

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```


Publisher _id as a Foreign Key

```
publisher = {  
  _id: "oreilly",  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
}
```

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  publisher_id: "oreilly"  
}
```

Book _id as a Foreign Key

```
publisher = {  
  name: "O'Reilly Media",  
  founded: "1980",  
  location: "CA"  
  books: [ "123456789", ... ]  
}
```

```
book = {  
  _id: "123456789",  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

Where Do You Put the Foreign Key?

Array of books inside of publisher

- Makes sense when **many** means a **handful** of items
- Useful when items have **bound** on potential **growth**

Reference to single publisher on books

- Useful when items have **unbounded growth** (unlimited # of books)

SQL doesn't give you a choice, no arrays

Another Example

One to Many Relations

Books and Patrons

Book can be checked out by one Patron at a time

Patrons can check out many books (but not 1000's)

Modeling Checkouts

```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader",  
  join_date: ISODate("2011-10-15"),  
  address: { ... }  
}
```

```
book = {  
  _id: "123456789",  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  ...  
}
```

Modeling Checkouts

```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader",  
  join_date: ISODate("2011-10-15"),  
  address: { ... },  
  checked_out: [  
    { _id: "123456789", checked_out: "2012-10-15" },  
    { _id: "987654321", checked_out: "2012-09-12" },  
    ...  
  ]  
}
```

Denormalization

Provides data locality

Modeling Checkouts: Denormalized

```
patron = {  
  _id: "joe",  
  name: "Joe Bookreader",  
  join_date: ISODate("2011-10-15"),  
  address: { ... },  
  checked_out: [  
    { _id: "123456789",  
      title: "MongoDB: The Definitive Guide",  
      authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
      checked_out: ISODate("2012-10-15")  
    },  
    { _id: "987654321"  
      title: "MongoDB: The Scaling Adventure",  
      ...  
    }, ...  
  ],  
}
```

Referencing vs. Embedding

- Embedding is a bit like pre-joined data
- Document-level ops are easy for server to handle
- Embed when the 'many' objects always appear with (i.e. viewed in the context of) their parent
- Reference when you need more flexibility

An Example

Single Table Inheritance

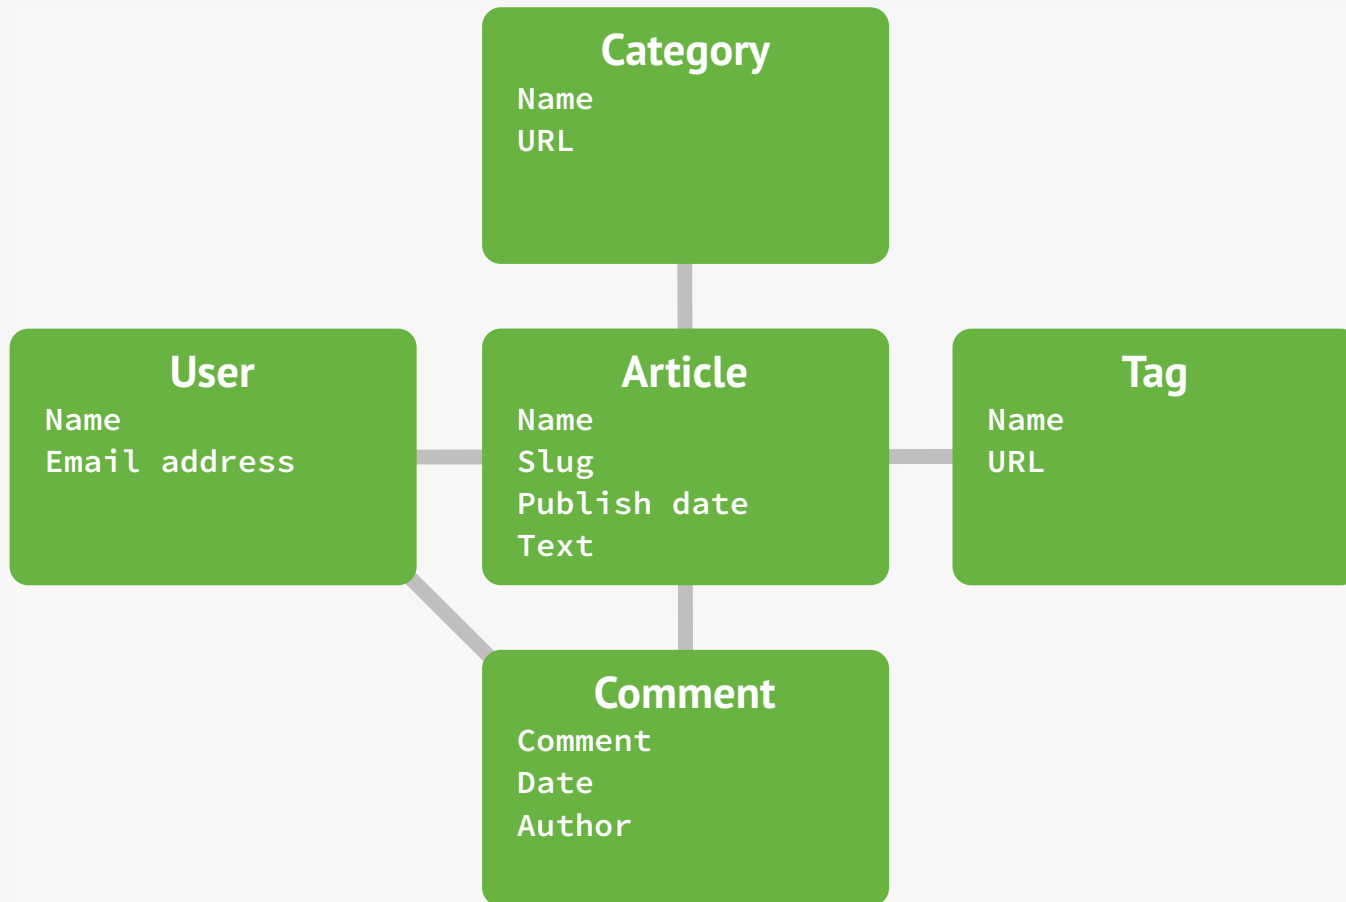
Single Table Inheritance

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  kind: "loanable",  
  locations: [ ... ],  
  pages: 216,  
  language: "English",  
  publisher: {  
    name: "O'Reilly Media",  
    founded: "1980",  
    location: "CA"  
  }  
}
```

An Example

Many to Many Relations

Relational Approach



Books and Authors

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors = [  
    { _id: "kchodorow", name: "K-Awesome" },  
    { _id: "mdirolf", name: "Batman Mike" },  
  ]  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
author = {  
  _id: "kchodorow",  
  name: "Kristina Chodorow",  
  hometown: "New York"  
}
```

Relation stored on both sides

```
book = {  
  _id: 123456789,  
  title: "MongoDB: The Definitive Guide",  
  authors = [ "kchodorow", "mdirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
author = {  
  _id: "kchodorow",  
  name: "Kristina Chodorow",  
  hometown: "Cincinnati",  
  books: [ 123456789, ... ]  
}
```


An Example

Trees

Parent Links

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  category: "MongoDB"  
}
```

```
category = { _id: MongoDB, parent: "Databases" }  
category = { _id: Databases, parent: "Programming" }
```

Child Links

```
book = {  
  _id: 123456789,  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English"  
}
```

```
category = { _id: MongoDB, children: [ 123456789, ... ] }
```

```
category = { _id: Databases, children: ["MongoDB", "Postgres"] }
```

```
category = { _id: Programming, children: ["DB", "Languages"] }
```

Modeling Trees

- Parent Links
 - Each node is stored as a document
 - Contains the id of the parent
- Child Links
 - Each node contains the id's of the children
 - Can support graphs (multiple parents / child)

Array of Ancestors

```
book = {  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  categories: ["Programming", "Databases", "MongoDB" ]  
}
```

```
book = {  
  title: "MySQL: The Definitive Guide",  
  authors: [ "Michael Kofler" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  parent: "MongoDB",  
  ancestors: [ "Programming", "Databases", "MongoDB" ]  
}
```

An Example

Queues

Book Document

```
book = {  
  _id: 123456789,  
  title: "MongoDB: The Definitive Guide",  
  authors: [ "Kristina Chodorow", "Mike Dirolf" ],  
  published_date: ISODate("2010-09-24"),  
  pages: 216,  
  language: "English",  
  available: 3  
}
```

```
db.books.findAndModify({  
  query: { _id: 123456789, available: { "$gt": 0 } },  
  update: { $inc: { available: -1 } }  
})
```



#ConferenceHashTag

Thank You

Speaker Name

Job Title, 10gen